

Dynamic Contract Checking for OCaml

– A Tutorial

Dana N. Xu
INRIA France
na.xu@inria.fr

May 17, 2011

Contents

1	Introduction	1
2	Contracts	3
2.1	Predicate contract	3
2.2	Dependent function contract	4
2.3	Pure functions in contracts	5
2.4	Specifying size properties	6
2.5	Higher order functions can be called in contracts	6
2.6	Higher order contracts	7
2.7	Dependent tuple contract	8
2.8	Local contract declaration	9
3	Divergent functions in contracts	11
4	Contract violations in contracts are reported	11
5	Modules	12
6	Command line instructions	13
7	Contract checking algorithm behind the scene	14
8	Future features	14

1 Introduction

A contract in a programming language is a formal and checkable interface specification that allows programmers to declare what a function assumes and what

a function guarantees. In an OCaml program, you can specify a contract for a top-level function with a keyword `contract`. Consider a simple example:

```
(* basic.ml *)
contract f = {x | x > 0} -> {r | r > x}
let f x = if x > 0 then x + 1
          else failwith "f requires a positive number"

let t1 = f 2
let t2 = f 0
```

The contract for `f` specifies that an input to `f` must be greater than 0 (indicated by `{x | x > 0}`) and `f` should guarantee that after taking a non-negative input, its output is greater than its input (indicated by `{r | r > x}`).

We can see two calls of `f` in `basic.ml`:

- `t1` calls `f` with 2, which satisfies `f`'s precondition;
- `t2` calls `f` with 0, which violates `f`'s precondition.

If we compile `basic.ml` and run it, we get:

```
Fatal error: exception
Contract_failure("basic.ml, line: 6, chars: 9-10,
Blame t2, because t2 fails f's precondition")
```

The error message gives the precise information where a contract violation occurs. It tells the programmer that in the file `basic.ml`, the call of `f` (at location (line 6, chars 9-10)) violates the contract and the function `t2` (i.e. the caller) is at fault.

Without the contract declaration, a call (`f 0`) gives the following run-time error message:

```
Fatal error: exception
Failure("f requires a positive number")
```

which does not tell which caller of `f` invokes this failure.

Consider another example:

```
(* basic.ml *)
contract g = {x | x > 0} -> {r | r > x}
let g x = x - 1

let t3 = g 2
```

We would get the following run-time error message:

```
Fatal error: exception
Contract_failure("basic.ml, line: 11, chars: 9-10,
Blame g, because g's postcondition fails")
```

It says that the call of `g` (at location (line 11, chars 9-10)) causes a contract violation and the function `g` itself is at fault. From the definition of `g`, we can see that the result of `g` is not greater than its input, thus, violates `g`'s postcondition.

From these examples, we can see that contracts give a way to reason about your program and help locating the function at fault precisely. This makes debugging more fun!

2 Contracts

Figure 1 gives the abstract syntax tree for contracts. We explain each of them in detail below. Note that the `->` is right associative and the `*` binds stronger than `->` without brackets.

$t ::=$	$\{x \mid p\}$	(* predicate contract *)
	$ \quad x:t_1 \rightarrow t_2$	(* dependent function contract *)
	$ \quad x:t_1 * t_2$	(* dependent tuple contract *)

Figure 1: Abstract Syntax Tree for Contracts

2.1 Predicate contract

We use set notation $\{x \mid p\}$ to mean $\forall x.p$ holds, where the p is a boolean-expression in OCaml. Currently, we allow *arbitrary pure functions not raising exceptions* to be called in the boolean expression p . For example, we may define a function `sorted`, which tests whether a list is sorted.

```
let rec sorted xs = match xs with
  | [] -> true
  | [x] -> true
  | x::y::l -> x <= y && sorted (y::l)
```

We may have a function `insert`, which inserts an item to a sorted list and returns a sorted list.

```
contract insert = {x | true} -> {y | sorted y} -> {r | sorted r}
let rec insert item xs = match xs with
  | [] -> [item]
  | x::l -> if item <= x then item::x::l
            else x::(insert item l)
```

We can now define a function `insertion_sort` which always returns a sorted list.

```
contract insertion_sort = {xs | true} -> {r | sorted r}
let rec insertion_sort xs = match xs with
  | [] -> []
  | x::l -> insert x (insertion_sort l)
```

As the function `insert` plays a key role, we would like to test it with `t1`.

```
let t1 = insert 5 [3;1;6]
```

We get this run-time error message:

```
Fatal error: exception
Contract_failure("insertionsort.ml, line: 19, chars: 9-15,
Blame t1, because t1 fails insert's precondition")
```

Clearly, it is `t1`'s fault not to provide a sorted list as `insert`'s second argument. For example, a programmer makes a careless mistake writing `>` instead of `<=`:

```
contract insert = {x | true} -> {y | sorted y} -> {r | sorted r}
let rec insert item xs = match xs with
| [] -> [item]
| x::l -> if item > x then item::x::l
          else x::(insert item l)
```

Testing it with `t2`:

```
let t1 = insert 5 [1;3;6]
```

We get a run-time error message:

```
Fatal error: exception
Contract_failure("insertionsort.ml, line: 19, chars: 9-15,
Blame insert, because insert's postcondition fails
or insert violates its parameter's contract")
```

which indicates that the careless mistake prevents `insert` producing a result that satisfies its postcondition. The second half of the message “`or insert violates its parameter's contract`” is useful for a function whose parameter is a function itself. The details can be found in Section 2.6.

2.2 Dependent function contract

Dependent function contracts allow us to specify a dependency between input and output in a function's postcondition. Its full notation is $x : t_1 \rightarrow t_2$ as shown in Figure 1 where x denotes an element in the precondition t_1 and can be used in the postcondition t_2 . A simple example of a dependent function contract is:

```
contract f = k:{x | x > 0} -> {r | r > k}
```

where `k` denotes an element in the precondition `{x | x > 0}` and we can refer to `k` in the scope of the postcondition `{r | ... }`. If `k` denotes an element of a predicate contract, we can omit the `k`, for example:

```
contract f = {x | x > 0} -> {r | r > x}
```

We allow such shorthand by assuming the x in the precondition $\{x \mid \dots\}$ scopes over the RHS of \rightarrow . However, if k denotes a (dependent) function contract, we cannot omit the $k:$. For example,

```
contract g = k:({x | x > 0} -> {y | y > x}) -> {z | z > 0}
          -> {r | r > k z}
let g f x = f x + 1
```

when we refer to the function contract, we have to use k in $\{r \mid r > k z\}$. Note that we *cannot* refer the x and the y outside the scope of the function contract, such as

```
contract g = k:({x | x > 0} -> {y | y > x}) -> {z | z = x}
          -> {r | r > k z}
```

2.3 Pure functions in contracts

Properties in contracts are all defined using pure expressions in OCaml itself. Logical operators are also pure functions. Note that the `&&` and `(||)` in the `stdlib` are lazily defined and can be used directly. If you want strict versions, you can have the following.

```
let (&&) x y = if x then y else false
let (||) x y = if x then true else y
let (==>) x y = if x then y else true
let not x = if x then false else true
```

Let us see more examples of what can be specified in contracts.

```
let null xs = match xs with
  | [] -> true
  | _ -> false

contract hd = {xs | not (null xs)} -> {r | true}
let hd xs = match xs with
  | [] -> failwith "hd"
  | x::_ -> x

contract tl = {xs | not (null xs)} -> {r | true}
let tl xs = match xs with
  | [] -> failwith "tl"
  | _::l -> l

contract risers = {xs | true}
          -> {r | not (null xs) ==> not (null r)}
let rec risers xs = match xs with
  | [] -> []
  | [x] -> [[x]]
```

```

| x::y::l -> let ss = risers (y::l) in
              if x <= y then (x::(hd ss))::(tl ss)
              else [x]::ss

```

We can see that preconditions of the functions `hd` and `tl` are always fulfilled.

2.4 Specifying size properties

In many cases, we want to specify size information of a function. For example:

```

contract length = {xs | true} -> {r | r >= 0}
let rec length xs = match xs with
  | [] -> 0
  | (x::l) -> 1 + length l

contract combine = {xs | true} -> {ys | length xs = length ys}
  -> {rs | length xs = length rs}
let rec combine xs ys = match (xs, ys) with
  | ([], []) -> []
  | (x::w, y::v) -> (x, y) :: combine w v
  | (_, _) -> failwith "requires arguments of equal length"

```

Function `combine` requires both arguments to have the same length and the length of the resulting list is the same as that of the first argument.

For another example:

```

let min x y = if x <= y then x else y

contract zip = {x | true} -> {y | true}
  -> {r | length rs = min (length x) (length y)}
let rec zip = xs ys =
  match (xs, ys) with
  | ([], _) | (_, []) -> []
  | (x::w, y::v) -> (x, y) :: zip w v

```

Function `zip` does the same as `combine` except that it stops combining two lists when the shorter list reaches its end.

2.5 Higher order functions can be called in contracts

We can define a higher-order function `filter` whose result is asserted with the help of another recursive higher-order function `for_all`.

```

(* val all : (a -> bool) -> [a] -> bool *)
let rec for_all f xs = match xs with
  | [] -> true
  | x::l -> f x && for_all f l

```

```

(* val filter : (a -> bool) -> [a] -> [a] *)
contract filter = {f | true} -> {x | true} -> {z | for_all f z}
let rec filter f xs = match xs with
  | [] -> []
  | x::l -> if f x then x :: filter f l
             else filter f l

```

2.6 Higher order contracts

A *higher order contract* allows t_1 in $x : t_1 \rightarrow t_2$ to be a (dependent) function contract itself. For example:

```

(* ho.ml *)
(* val f1 : (int -> int) -> int *)
contract f1 = ({x | x > 0} -> {y | y >= x}) -> {z | z > 0}
let f1 g = (g 0) - 5

let t1 = f1 (fun x -> x + 1)

```

In this case, we get run-time error message:

```

Fatal error: exception
Contract_failure("ho.ml, line: 5, chars: 9-11,
Blame f1, because f1's postcondition fails or
f1 violates its parameter's contract")

```

It is because `f1` does not use `g` in a way that satisfies `g`'s precondition. But this is detected at the call site of `f1`, i.e. (line 6, chars 9-11).

For another example:

```

contract f2 = ({x | x > 0} -> {y | y >= x}) -> {z | z > 0}
let f2 g = (g 1) - 5

let t2 = f2 (fun x -> x + 1)

```

We get the following error message:

```

Fatal error: exception
Contract_failure("ho.ml, line: 9, chars: 9-11,
Blame f2, because f2's postcondition fails")

```

We can see that the argument of `f2`, `(fun x -> x + 1)`, satisfies `f2`'s precondition. From the contract of `f2`, we can see that $x > 0$ and $y \geq x$, so $y > 0$, that is, $(g\ 1) > 0$. However, $(g\ 1) - 5$ is not greater than 0, thus, `f2` fails its postcondition $\{z \mid z > 0\}$. That is why the error message blames `f2`. The location (line 9, chars 9-11) refers to the call site of `f2`.

Moreover, if we have:

```

let t3 = f2 (fun x -> x - 1)

```

we have error message:

```
Fatal error: exception
Contract_failure("ho.ml, line: 11, chars: 9-11,
Blame t3, because t3 fails f2's precondition")
```

The location (line 11, chars 9-11) is the call site of `f2`. Since `t3` calls `f2` with a function that fails `f2`'s precondition, the function `t3` is blamed.

2.7 Dependent tuple contract

Currently, we allow tuple contract of arbitrary length. For example, we may have:

```
(* tuple.ml *)
let fst (a,b) = a
let snd (a,b) = b
(* val f : int array -> int -> int *)
contract f = k:({arr | true} * {i | 0 <= i && i < Array.length arr})
-> {r | r = snd k}
let f (arr, i) = arr.(i)
let a = Array.init 5 (fun x -> x)
let t1 = f (a, 2)
let t2 = f (a, 8)
```

Function `f` takes a tuple of an array `arr` and an index `i` and access the `i`th element in the array. The contract of `f` requires the index to be within the bounds of the array, i.e. greater or equal to 0 and less than the length of the array. It also specifies that the result of `f` should be the same as the second component of the argument (i.e. the index). We can see that the call in `t1` satisfies `f`'s precondition while the call in `t2` violates `f`'s precondition and gives the following run-time error message:

```
Fatal error: exception
Contract_failure("tuple.ml, line: 9, chars: 9-10,
Blame t2, because t2 fails f's precondition")
```

Given a dependent tuple contract $x:t_1 * t_2$, if t_1 denotes a predicate contract, we can use shorthand, for example,

```
{arr | true} * {i | 0 <= i && i < Array.length arr}
```

by assuming the variable `arr` scopes over the second component. However, if t_1 is a (dependent) function contract, we cannot omit the x in $x:t_1 * t_2$. For example:

```
k:({x | x > 0} -> {y | y > x}) * {z | z > k 0}
```


we cannot omit the `k` if we want to use it in `{z | ... }`.

Moreover, in the contract of `f`, we can see that the variable `arr` and `i` do not scope over the RHS of `->`. In order to refer to the subcomponents of the argument, we have to define projection function that projects on the subcomponent such as `fst` and `snd` and use them instead.

Consider another example:

```
contract f1 = {x | x > 0} * {y | y > 0} -> {z | z > 0} ;;
let f1 (a,b) = a
let t3 = f1 (5, -1)
```

Although the second element of the tuple is not used in the definition of `fstPos`. Due to strict evaluation, a call `fst (5, -1)` fails `fst`'s precondition.

Recall the syntax in Figure 1, we can also give `fstPos` the following contract.

```
contract fstPos = {x | fst x > 0 && snd x > fst x}
               -> {r | r = fst x}
let fstPos (a,b) = a
```

Here, we use a predicate contract for the parameter of `fstPos` and requires the second component to be greater than the first.

If we are only interested in the first element of the tuple, we can write:

```
contract fstPos = {x | fst x > 0} -> {r | r = fst x}
```

A call `fstPos (5, -1)` satisfies `fstPos`'s precondition. However, a call

```
fstPos (5, failwith "fstPos: 2nd component")
```

gives error message:

```
Fatal error: exception Failure("fstPos: 2nd component")
```

It is because the contract `{x | fst x > 0}` requires that the first component of `x` to be greater than 0 and the second component to be any total expression. A *total* expression refers to an expression that neither diverges nor throws an exception.

2.8 Local contract declaration

It is possible to give a contract to a local let-binding. For example:

```
contract bubblesort = {xs | true} -> {r | sorted r}
let rec bubblesort xs =
  let rec bsorthelper sat {xs | true} ->
    {r | (not (snd r)) ==> (sorted (fst r))}
  = function
  | [] -> ([], false)
  | [a] -> ([a], false)
  | x::xs -> let (y::ys, changed) = bsorthelper xs in
```

```

        if x <= y then (x :: (y::ys), changed)
        else (y :: (x::ys), true)
    in
    let (result, changed) = bsorthelper xs in
    if changed then bubblesort result
    else result

```

```
let test_bubblesort = bubblesort [2,1,5]
```

We give a contract to the local function `bsorthelper` in a format

$$\langle \text{function name} \rangle \text{ sat } \langle \text{contract} \rangle = \langle \text{function body} \rangle$$

which is analogic to the local function type declaration

$$\langle \text{function name} \rangle : \langle \text{type} \rangle = \langle \text{function body} \rangle$$

In the above example, the result of `bubblesort` is sorted and thus satisfies its postcondition. Let us see another example:

```

contract f = {x | x >= 0} -> {y | y > x}
let f v = let rec h sat {x | x > 5} -> {y | y > x}
          = fun x -> x + v
          and k sat {x | x > 3} -> {y | y > 3}
          = fun x -> h x + 1 in
          k v + h 6

```

```

let t1 = f 0
let t2 = f 4

```

The contract of the local function `h` requires its result to be greater than its input. The call `(f 0)` in the definition of `t1` produces `x+0` which is not greater than `x`. Thus, the postcondition of `h` is not satisfied and we get run-time error message:

```

Fatal error: exception
Contract_failure("local.ml, line: 6, chars: 17-18,
Blame h, because h's postcondition fails or
h violates its parameter's contract")

```

where the location (line 6, chars 17-18) refers to the call `(h 6)`.

The local function `k` requires its input to be greater than 3, the call `(f 4)` in the definition of `t2` satisfies this precondition. However, the function `k` calls the local function `h` with its argument and fails `h`'s precondition `{x | x > 5}`. Thus, the call `h x` (line 5, chars 26-27), in the definition of `k`, violates `h`'s precondition and the function `k` is to blame. So we get the following run-time error message:

```

Fatal error: exception
Contract_failure("local.ml, line: 5, chars: 26-27,
Blame k, because k fails h's precondition")

```

3 Divergent functions in contracts

We only have partial correctness for our contract framework. That is, a divergent expression satisfies any contract and a divergent expression is the only expression that satisfies $\{x \mid p\}$ where p evaluates to **false**.

Given

```
(* loop.ml *)
contract loop = {x | x > 0} -> {y | false}
let rec loop x = loop x
let t1 = loop 0
let t2 = loop 5
```

The call `(loop 0)` invokes error message:

```
Fatal error: exception
Contract_failure("loop.ml, line: 4, chars: 9-13,
Blame t1, because t1 fails loop's precondition")
```

while a call `(loop 5)` diverges (i.e. causes stack overflow).

4 Contract violations in contracts are reported

As we allow arbitrary pure functions to be called in contracts, it is possible that a call to a function throws an exception in contracts. For example:

```
(* c.ml *)
contract f1 = {x | true} -> {y | hd x > y}
let f1 x = 5

let t1 = f1 []
```

As we gave `hd` a contract, we can get run-time error message:

```
Fatal error: exception
Contract_failure("cntr.ml, line: 16, chars: 33-34,
Blame f1, because f1 fails hd's precondition")
```

The location (line 16, chars 33-34) refers to the call site of `hd` in the contract $\{y \mid \text{hd } x > y\}$. Here, "blame: f1" means blaming `f1`'s contract.

For another example:

```
contract f2 = k:({x | x > 0} -> {y | y > x}) -> {z | k 0 < z}
let f2 g = g 5

let t2 = f2 (fun x -> x + 1)
```

We have error message:

```
Fatal error: exception
Contract_failure("cntr.ml, line: 7, chars: 53-54,
Blame f2, because f2 fails k's precondition")
```

We blame the location of the call (`k 0`), which is in the contract of `f2`. Note, in this case, the `f2` in `f2 fails k's precondition` refers to the contract of `f2`.

To use a partial function in contracts that may throw an exception, we recommend users to give sufficient precondition for the partial function used. For example:

```
contract f3 = {xs | not (null xs)} -> {z | z > hd xs}
let f3 xs = match xs with
  | [] -> failwith "f3"
  | x::xs -> x + 1

let t3 = f3 []
```

We get error message:

```
Fatal error: exception
Contract_failure("cntr.ml, line: 18, chars: 9-11,
Blame t3, because t3 fails f3's precondition")
```

We can see that the precondition of `hd` is always satisfied when the call `hd xs` is reached.

5 Modules

We allow a contract to be defined directly next to its corresponding function in `.ml` files with a keyword `contract` mainly because programmers often want to specify contracts for functions that are not exported as well. Exported contracts are in `.mli` files where they should be syntactically the same as those in `.ml` files and declared after their type declaration. (That is, if a contract is specified for an exported function, programmers should write this contract in both `.ml` and `.mli` files and make sure that they are syntactically the same.) For example:

```
(* basic.mli *)
val f: int -> int
contract f = {x | x >= 0} -> {y | y > x}
```

It is possible not to write a `.mli` file. In this case, all contracts declared in the `.ml` file are presumed to be exported.

Similar ideas apply to nested modules: exported contracts are written in a module signature; if no module signature is given, we assume all contracts declared in the module are exported.

6 Command line instructions

Source code can be obtained by:

```
svn checkout https://yquem.inria.fr/caml/svn/ocaml/branches/contracts
```

which is based on the source code of ocamlc-3.11.2. To build ocamlc from source, do:

```
./configure
make world
```

Small examples can be downloaded at:

```
http://gallium.inria.fr/~naxu/research/testcontracts.tar.bz2
```

I put the folder `testcontracts` and the folder `contracts` side by side in the same directory. To release .tar.bz2, do:

```
tar -jxvf testcontracts.tar.bz2
```

To test one of the examples, e.g. `basic.ml`, you can do:

```
make basic
ocamlrun ./basic
```

To compile all examples:

```
make all
```

If you have installed the newly built ocamlrun and ocamlc and you do not want to use a Makefile, you can compile your own example as usual:

```
ocamlrun ocamlc -I stdlib -c basic.ml
ocamlrun ocamlc -o basic -I stdlib basic.cmo
ocamlrun ./basic
```

If you want to install the newly built ocamlc locally, you need do:

```
make distclean
./configure --prefix='pwd'/_install
make world
make install
```

This will install everything at local directory `./_install`. For example, you can see ocamlc, ocamlrun at `_install/bin/` and the standard library is at `_install/lib/ocaml/`. You need to give the full path to invoke them.

If you want to switch off the contract checking, you can do:

```
ocamlrun ocamlc -I stdlib -nocontract -c basic.ml
```

You can look up the flag `-nocontract` with `ocamlc -help`.

7 Contract checking algorithm behind the scene

The beauty of contract checking is that *it is simple!* Given

```
contract f = t
let f = e
```

where the t refers to the contract of the function \mathbf{f} and the e refers to the definition of \mathbf{f} . The contract checking algorithm is:

$$\begin{aligned}
e \underset{l_2}{\overset{l_1}{\boxtimes}} \{x \mid p\} &= \text{let } x = e \text{ in if } p \text{ then } x \text{ else blame } l_1 \\
e \underset{l_2}{\overset{l_1}{\boxtimes}} x:t_1 \rightarrow t_2 &= \text{fun } v \rightarrow (e \underset{l_1}{\overset{l_2}{\boxtimes}} t_1) \underset{l_2}{\overset{l_1}{\boxtimes}} t_2[(v \underset{l_1}{\overset{l_2}{\boxtimes}} t_1)/x] \\
e \underset{l_2}{\overset{l_1}{\boxtimes}} x:t_1 * t_2 &= \text{match } e \text{ with} \\
&\quad (x_1, x_2) \rightarrow (x_1 \underset{l_2}{\overset{l_1}{\boxtimes}} t_1, x_2 \underset{l_2}{\overset{l_1}{\boxtimes}} t_2[(x_1 \underset{l_1}{\overset{l_2}{\boxtimes}} t_1)/x])
\end{aligned}$$

The operator $\underset{l_2}{\overset{l_1}{\boxtimes}}$ is called a wrapper, which wraps an expression with its contract.

If the contract is satisfied, the original expression is returned; otherwise, a contract violation is reported. Here the l_i contains information about source location and function name of a callee (or a caller). The algorithm makes blaming of a callee or a caller precise. For more information related to contracts, you can read [1, 2].

8 Future features

More features we are considering include: conjunctive and disjunctive contracts, invariants for data types, polymorphic contracts, contract synonyms, etc.

Acknowledgement I would like to thank Xavier Leroy and Nicolas Pouillard for explaining the internal structure of the OCaml compiler.

References

- [1] Robert Bruce Findler and Matthias Felleisen. Contracts for higher-order functions. In *ICFP '02: Proceedings of the seventh ACM SIGPLAN international conference on Functional programming*, pages 48–59, New York, NY, USA, 2002. ACM Press.
- [2] Dana N. Xu, Simon L. Peyton Jones, and Koen Claessen. Static contract checking for haskell. In Zhong Shao and Benjamin C. Pierce, editors, *POPL*, pages 41–52. ACM, 2009.